

Prerequisites: Setup the environment

I have a ubuntu box and already had virtualbox installed. So decided to use vagrant to setup the environment. Instructions:

<https://www.vagrantup.com/downloads.html>

Screenshot:

```
gmulchan@gmulchan-XPS-8900:~/Downloads$ sudo dpkg -i vagrant_2.1.2_x86_64.deb
[sudo] password for gmulchan:
Selecting previously unselected package vagrant.
(Reading database ... 426331 files and directories currently installed.)
Preparing to unpack vagrant_2.1.2_x86_64.deb ...
Unpacking vagrant (1:2.1.2) ...
Setting up vagrant (1:2.1.2) ...
```

Start vagrant:

```
gmulchan@gmulchan-XPS-8900:~/Downloads$ cd
gmulchan@gmulchan-XPS-8900:~$ mkdir datadog
gmulchan@gmulchan-XPS-8900:~$ cd datadog/
gmulchan@gmulchan-XPS-8900:~/datadog$ ls
gmulchan@gmulchan-XPS-8900:~/datadog$ vagrant init
A 'Vagrantfile' has been placed in this directory. You are now
ready to 'vagrant up' your first virtual environment! Please read
the comments in the Vagrantfile as well as documentation on
'vagrantup.com' for more information on using Vagrant.
gmulchan@gmulchan-XPS-8900:~/datadog$ ls
Vagrantfile
gmulchan@gmulchan-XPS-8900:~/datadog$ vi Vagrantfile
gmulchan@gmulchan-XPS-8900:~/datadog$ cp Vagrantfile Vagrantfile.backup
gmulchan@gmulchan-XPS-8900:~/datadog$ rm Vagrantfile
gmulchan@gmulchan-XPS-8900:~/datadog$ vi Vagrantfile
gmulchan@gmulchan-XPS-8900:~/datadog$
gmulchan@gmulchan-XPS-8900:~/datadog$
gmulchan@gmulchan-XPS-8900:~/datadog$
gmulchan@gmulchan-XPS-8900:~/datadog$ vagrant up
Bringing machine 'default' up with 'virtualbox' provider...
==> default: Box 'gbarbieru/xenial' could not be found. Attempting to find and install...
    default: Box Provider: virtualbox
    default: Box Version: >= 0
==> default: Loading metadata for box 'gbarbieru/xenial'
    default: URL: https://vagrantcloud.com/gbarbieru/xenial
==> default: Adding box 'gbarbieru/xenial' (v0.0.6) for provider: virtualbox
    default: Downloading: https://vagrantcloud.com/gbarbieru/boxes/xenial/versions/0.0.6/providers/virtualbox.box
==> default: Successfully added box 'gbarbieru/xenial' (v0.0.6) for 'virtualbox'!
==> default: Importing base box 'gbarbieru/xenial'...
==> default: Matching MAC address for NAT networking...
==> default: Checking if box 'gbarbieru/xenial' is up to date...
==> default: Setting the name of the VM: datadog_default_1530834184513_98455
Vagrant is currently configured to create VirtualBox synced folders with
```

Login to Vagrant, create account at datadog and Install Datadog agent. Datadog agent already has the right keys so that it can report to the correct account, the default hostname looks good.

Screenshot of installing agent:

```
vagrant@vagrant: /etc/datadog-agent
gmulcha... x vagrant... x vagrant... x vagrant... x gmulcha... x
==> default: Configuring and enabling network interfaces...
==> default: Mounting shared folders...
default: /vagrant => /home/gmulchan/datadog
gmulchan@gmulchan-XPS-8900:~/datadog$ vagrant ssh
Welcome to Ubuntu 16.04 LTS (GNU/Linux 4.4.0-21-generic x86_64)

 * Documentation:  https://help.ubuntu.com/

196 packages can be updated.
108 updates are security updates.

Last login: Thu Apr 21 05:00:28 2016
vagrant@vagrant:~$ ls
vagrant@vagrant:~$ pwd
/home/vagrant
vagrant@vagrant:~$ ls
vagrant@vagrant:~$ DD_API_KEY=a32ed5c86d3cb70caac39a17d0800f1b bash -c
"$(<curl -L https://raw.githubusercontent.com/DataDog/datadog-agent/master/cmd/agent/install_script.sh)"
  % Total    % Received % Xferd  Average Speed   Time    Time     Time
  Current                                 Dload  Upload  Total   Spent    Left
Speed
100 10240  100 10240    0     0  22004      0 --:--:-- --:--:-- --:--:--
- 21974

* Installing apt-transport-https

Hit:1 http://us.archive.ubuntu.com/ubuntu xenial InRelease
Hit:2 http://us.archive.ubuntu.com/ubuntu xenial-updates InRelease
Hit:3 http://security.ubuntu.com/ubuntu xenial-security InRelease
Hit:4 http://us.archive.ubuntu.com/ubuntu xenial-backports InRelease
Reading package lists...
Reading package lists...
Building dependency tree...
Reading state information...
The following NEW packages will be installed:
  apt-transport-https
0 upgraded, 1 newly installed, 0 to remove and 115 not upgraded.
Need to get 26.1 kB of archives.
After this operation, 215 kB of additional disk space will be used.
Get:1 http://us.archive.ubuntu.com/ubuntu xenial-updates/main amd64 apt-transport-https amd64 1.2.27 [26.1 kB]
Fetched 26.1 kB in 0s (63.5 kB/s)
Selecting previously unselected package apt-transport-https.
(Reading database ... 25535 files and directories currently installed.)
Preparing to unpack .../apt-transport-https_1.2.27_amd64.deb ...
Unpacking apt-transport-https (1.2.27) ...
Setting up apt-transport-https (1.2.27) ...
Reading package lists...
```


Collecting metrics:

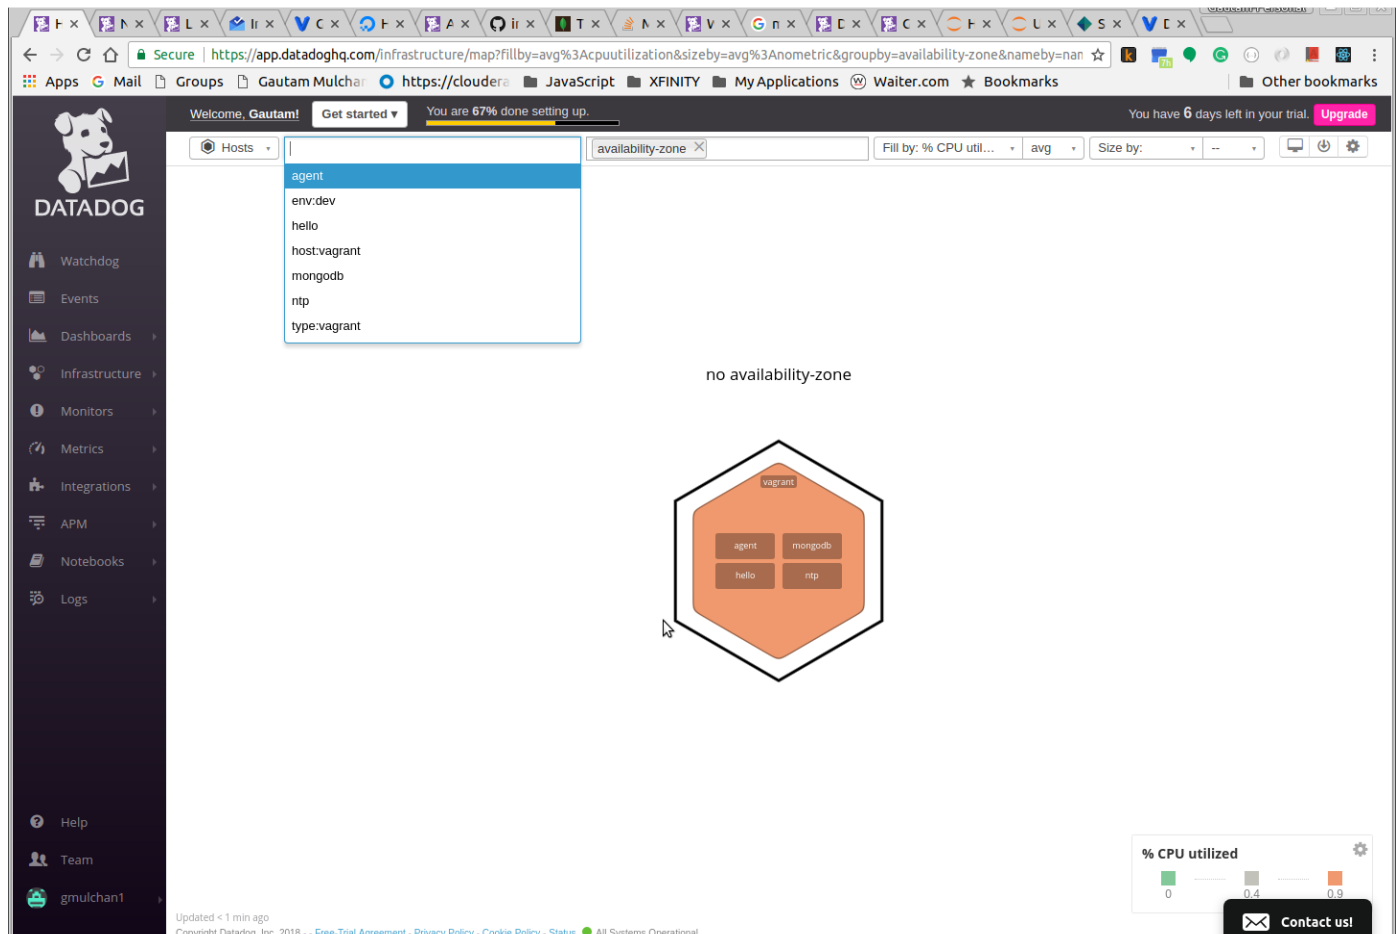
- Add tags in the Agent config file and show us a screenshot of your host and its tags on the Host Map page in Datadog.

Assigning tags to components allows analytics on the reporting. One can slice and dice the reports, aggregate and group metrics based on tags (for example, dev/stage/prod envs). This is a concept from business analytics increasingly applied to operations and service management. Objective is to better understand metrics as a group of entities where the entities can be dynamically grouped using tags as a mechanism. Used mechanism of adding tags to the datadog.yaml file. Screenshot below:

https://docs.datadoghq.com/getting_started/tagging/assigning_tags/

```
# Set the host's tags (optional)
# tags:
#   - mytag
#   - env:prod
#   - role:database
tags: env:dev, type:vagrant
```

And the tag shows up in datadog UI making it easier to search:



In this case there is only one host so benefits of tagging is only for searching. With larger amounts of data, tagging helps in reporting and analytics.

- Install a database on your machine (MongoDB, MySQL, or PostgreSQL) and then install the respective Datadog integration for that database.

Installed mongodb on the vagrant host.

```
vagrant@vagrant: /etc/datadog-agent/conf.d
gmulcha... x vagrant... x vagrant... x vagrant... x gmulcha... x
gmulchan@gmulchan-XPS-8900:~/datadog$ vagrant ssh
Welcome to Ubuntu 16.04 LTS (GNU/Linux 4.4.0-21-generic x86_64)

 * Documentation:  https://help.ubuntu.com/

120 packages can be updated.
6 updates are security updates.

*** System restart required ***
Last login: Thu Jul  5 19:43:54 2018 from 10.0.2.2
vagrant@vagrant:~$ ls
ddagent-install.log
vagrant@vagrant:~$ sudo apt-key adv --keyserver hkp://keyserver.ubuntu.
com:80 --recv EA312927
Executing: /tmp/tmp.xCz1lSdYHk/gpg.1.sh --keyserver
hkp://keyserver.ubuntu.com:80
--recv
EA312927
gpg: requesting key EA312927 from hkp server keyserver.ubuntu.com
gpg: key EA312927: public key "MongoDB 3.2 Release Signing Key <packagi
ng@mongodb.com>" imported
gpg: Total number processed: 1
gpg:          imported: 1 (RSA: 1)
vagrant@vagrant:~$ echo "deb http://repo.mongodb.org/apt/ubuntu xenial/
mongodb-org/3.2 multiverse" | sudo tee /etc/apt/sources.list.d/mongodb-
org-3.2.list
deb http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2 multivers
e
vagrant@vagrant:~$ sudo apt-get update
Ign:1 http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2 InRelea
se
Hit:2 http://us.archive.ubuntu.com/ubuntu xenial InRelease

Get:3 http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2 Release
[3,462 B]
Get:4 http://us.archive.ubuntu.com/ubuntu xenial-updates InRelease [109
kB]
Get:5 http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2 Release
.gpg [801 B]
Get:6 http://security.ubuntu.com/ubuntu xenial-security InRelease [107
kB]
Get:7 http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2/multive
rse amd64 Packages [10.1 kB]
Get:8 http://us.archive.ubuntu.com/ubuntu xenial-backports InRelease [1
07 kB]
Get:9 http://us.archive.ubuntu.com/ubuntu xenial-updates/main amd64 Pac
kages [804 kB]
Get:10 http://security.ubuntu.com/ubuntu xenial-security/main amd64 Pac
kages [519 kB]
Get:11 http://us.archive.ubuntu.com/ubuntu xenial-updates/main i386 Pac
kages [734 kB]
```



```
vagrant@vagrant:~$ sudo apt-get install -y mongodb-org
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  mongodb-org-mongos mongodb-org-server mongodb-org-shell mongodb-org-tools
The following NEW packages will be installed:
  mongodb-org mongodb-org-mongos mongodb-org-server mongodb-org-shell
  mongodb-org-tools
0 upgraded, 5 newly installed, 0 to remove and 115 not upgraded.
Need to get 51.7 MB of archives.
After this operation, 214 MB of additional disk space will be used.
Get:1 http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2/multiverse amd64 mongodb-org-shell amd64 3.2.20 [5,277 kB]
Get:2 http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2/multiverse amd64 mongodb-org-server amd64 3.2.20 [10.0 MB]
Get:3 http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2/multiverse amd64 mongodb-org-mongos amd64 3.2.20 [4,677 kB]
Get:4 http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2/multiverse amd64 mongodb-org-tools amd64 3.2.20 [31.8 MB]
Get:5 http://repo.mongodb.org/apt/ubuntu xenial/mongodb-org/3.2/multiverse amd64 mongodb-org amd64 3.2.20 [3,550 B]
Fetched 51.7 MB in 42s (1,207 kB/s)
```

```
Done.
Setting up mongodb-org-mongos (3.2.20) ...
Setting up mongodb-org-tools (3.2.20) ...
Setting up mongodb-org (3.2.20) ...
vagrant@vagrant:~$ sudo systemctl start mongod
vagrant@vagrant:~$ sudo systemctl status mongod
● mongod.service - High-performance, schema-free document-oriented data
base
   Loaded: loaded (/lib/systemd/system/mongod.service; disabled; vendor
 preset:
   Active: active (running) since Tue 2018-07-10 19:38:36 EDT; 11s ago
     Docs: https://docs.mongodb.org/manual
    Main PID: 17366 (mongod)
      CGroup: /system.slice/mongod.service
              └─17366 /usr/bin/mongod --quiet --config /etc/mongod.conf

Jul 10 19:38:36 vagrant systemd[1]: Started High-performance, schema-fr
ee docume

vagrant@vagrant:~$
vagrant@vagrant:~$
vagrant@vagrant:~$ sudo systemctl enable mongod
Created symlink from /etc/systemd/system/multi-user.target.wants/mongod
.service to /lib/systemd/system/mongod.service.
vagrant@vagrant:~$ cd /etc/
```

Enabled access from datadog agent to query mongodb to collect metrics:

```

vagrant@vagrant:/etc$ vi mongod.conf
vagrant@vagrant:/etc$ mongo
MongoDB shell version: 3.2.20
connecting to: test
Welcome to the MongoDB shell.
For interactive help, type "help".
For more comprehensive documentation, see
    http://docs.mongodb.org/
Questions? Try the support group
    http://groups.google.com/group/mongodb-user
Server has startup warnings:
2018-07-10T19:38:36.944-0400 I CONTROL [initandlisten]
2018-07-10T19:38:36.944-0400 I CONTROL [initandlisten] ** WARNING: /sys
s/kernel/mm/transparent_hugepage/enabled is 'always'.
2018-07-10T19:38:36.944-0400 I CONTROL [initandlisten] **          We su
ggest setting it to 'never'
2018-07-10T19:38:36.944-0400 I CONTROL [initandlisten]
2018-07-10T19:38:36.945-0400 I CONTROL [initandlisten] ** WARNING: /sys
s/kernel/mm/transparent_hugepage/defrag is 'always'.
2018-07-10T19:38:36.945-0400 I CONTROL [initandlisten] **          We su
ggest setting it to 'never'
2018-07-10T19:38:36.945-0400 I CONTROL [initandlisten]
> use admin
switched to db admin
> db.addUser("datadog", "datadog", true)
2018-07-10T19:55:57.153-0400 E QUERY [thread1] TypeError: db.addUser
is not a function :
@ (shell):1:1

> db.createUser({"user": "datadog", "pwd": "datadog", "roles": [{role: '
read', db: 'admin' }, {role: 'clusterMonitor' , db: 'admin'}, {role: 'r
ead', db: 'local' } ] } )
Successfully added user: {
  "user" : "datadog",
  "roles" : [
    {
      "role" : "read",
      "db" : "admin"
    },
    {
      "role" : "clusterMonitor",
      "db" : "admin"
    },
    {
      "role" : "read",
      "db" : "local"
    }
  ]
}

```



```

vagrant@vagrant:/var/log/mongodb$
vagrant@vagrant:/var/log/mongodb$
vagrant@vagrant:/var/log/mongodb$
vagrant@vagrant:/var/log/mongodb$ sudo systemctl status mongod
● mongod.service - High-performance, schema-free document-oriented data
   Loaded: loaded (/lib/systemd/system/mongod.service; enabled; vendor
   Active: active (running) since Tue 2018-07-10 19:38:36 EDT; 18h ago
     Docs: https://docs.mongodb.org/manual
   Main PID: 17366 (mongod)
    CGroup: /system.slice/mongod.service
            └─17366 /usr/bin/mongod --quiet --config /etc/mongod.conf

Jul 10 19:38:36 vagrant systemd[1]: Started High-performance, schema-fr

```

Change the config in the mongodb config file to monitor the local mongodb instance. Also enabled logs monitoring.

```

vagrant@vagrant: /etc/datadog-agent/conf.d/mongo.d
init_config:
instances:
  - server: mongodb://datadog:datadog@localhost:27017/admin
    additional_metrics:
      - collection # collect metrics for each collection
      - metrics.commands
      - tcmalloc
      - top

logs:
  - type: file
    path: /var/log/mongodb/mongod.log
    service: mongoDB
    source: mongodb

~
~
1,1 All

```

Restarted agent after enabling mongodb monitoring. Hopefully should not need to restart agent after enabling/disabling any component. Agent should pick up the changes in config automatically. This needs to be done for every step.

```

vagrant@vagrant:/etc/datadog-agent/conf.d/mongo.d$
vagrant@vagrant:/etc/datadog-agent/conf.d/mongo.d$
vagrant@vagrant:/etc/datadog-agent/conf.d/mongo.d$ sudo systemctl stop
datadog-agent
vagrant@vagrant:/etc/datadog-agent/conf.d/mongo.d$ sudo systemctl start
datadog-agent
vagrant@vagrant:/etc/datadog-agent/conf.d/mongo.d$ sudo service datadog
-agent status
● datadog-agent.service - "Datadog Agent"
   Loaded: loaded (/lib/systemd/system/datadog-agent.service; enabled;
   vendor pr
   Active: active (running) since Wed 2018-07-11 13:54:48 EDT; 9min ago
   Main PID: 25513 (agent)
   CGroup: /system.slice/datadog-agent.service
           └─25513 /opt/datadog-agent/bin/agent/agent start -p /opt/dat
   adog-agent

```

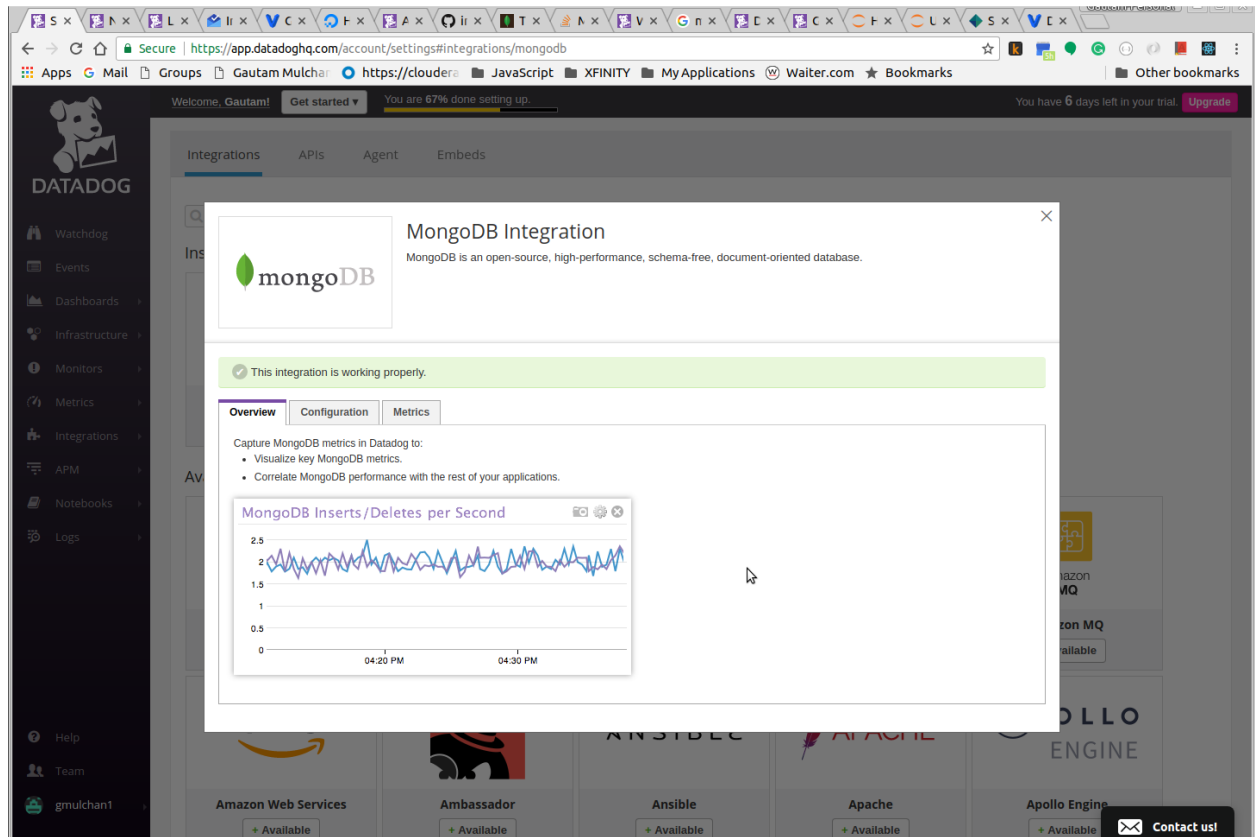
Mongodb metrics not showing up in datadog UI, looked at datadog agent log and saw an error:

```

2018-07-10 20:00:22 EDT | INFO | (collector.go:52 in NewCollector) | Em
bedding Python 2.7.14 (default, Jul  4 2018, 16:21:37) [GCC 4.7.2]
2018-07-10 20:00:22 EDT | INFO | (file.go:69 in Collect) | File Configu
ration Provider: searching for configuration files at: /etc/datadog-age
nt/conf.d
2018-07-10 20:00:22 EDT | WARN | (file.go:179 in collectEntry) | /etc/d
atadog-agent/conf.d/mongo.d/conf.yaml is not a valid config file: Confi
guration file contains no valid instances
2018-07-10 20:00:22 EDT | INFO | (file.go:69 in Collect) | File Configu
ration Provider: searching for configuration files at: /opt/datadog-age
nt/bin/agent/dist/conf.d

```

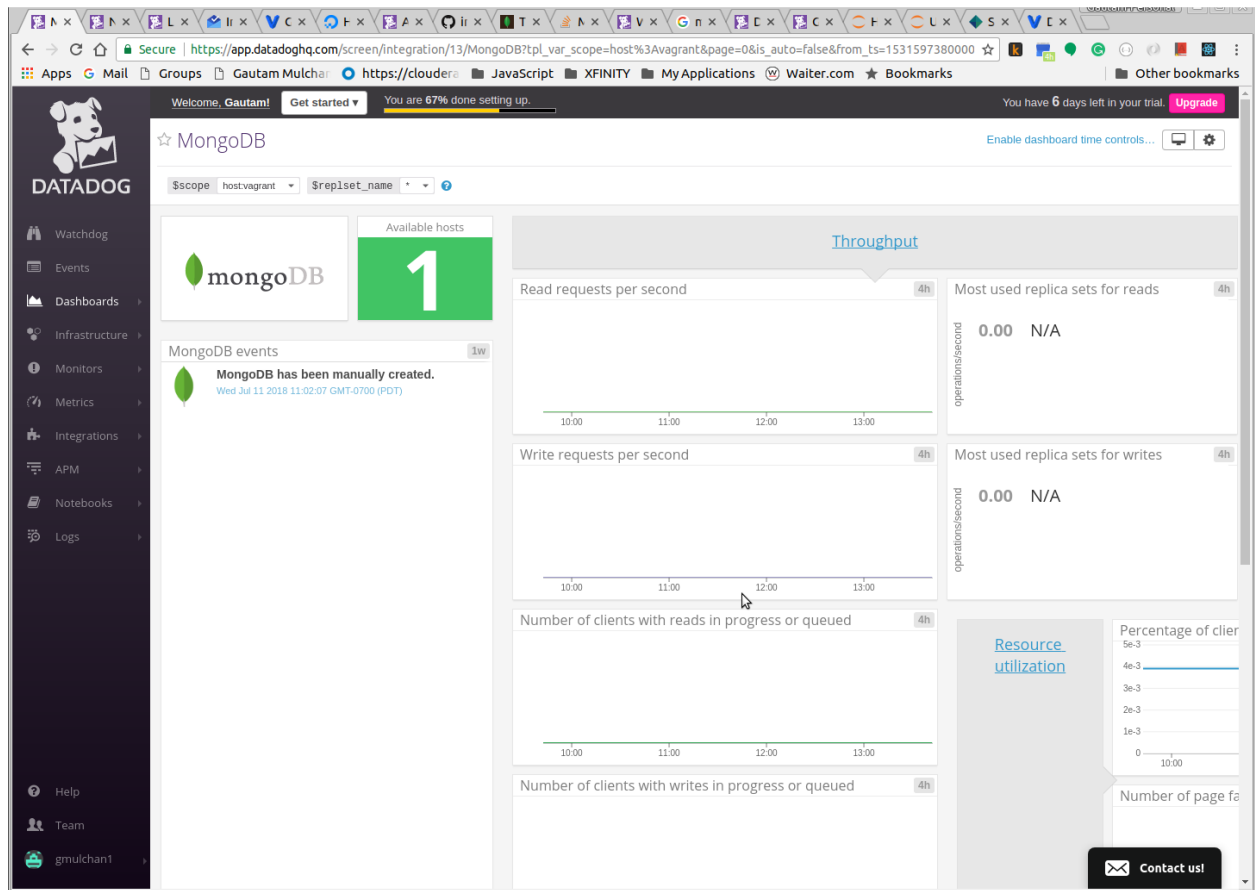
Turned out a typo in the yaml file. In meantime also enabled the mongodb integration.



It is unclear what integrations (in the UI) do. The documentation lists the over 200 integrations but does not explain what specifically is the purpose of integrations and how exactly the integration helps from a use case perspective.

<https://docs.datadoghq.com/integrations/>

Here is mongodb dashboard:



Overall, very easy to setup monitoring for various infrastructure components with just a small config change in the agent and everything else in UI.

- Create a custom Agent check that submits a metric named my_metric with a random value between 0 and 1000.

Same comment about Agent Checks as for the integrations, will be great to document the purpose and how it helps in monitoring.

https://docs.datadoghq.com/developers/agent_checks/

```
vagrant@vagrant: /etc/datadog-agent/conf.d
init_config:
  min_collection_interval: 45

instances:
  - min_collection_interval: 45

~
~
~
~
1,1 All
```

```

vagrant@vagrant: /etc/datadog-agent/checks.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 supervisord.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 marathon.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 kubelet.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 fluentd.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 activemq.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 varnish.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 memory.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 kube_dns.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 riakcs.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 mesos_master.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 hdfs_namenode.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 gunicorn.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 kafka_consumer.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 istio.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 haproxy.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 envoy.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 uptime.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 snmp.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 mysql.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 gitlab.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 couchbase.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 zk.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 tcp_check.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 nfsstat.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 kubernetes_apiserver.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 system_swap.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 network.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 ecs_fargate.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 docker.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 activemq_xml.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 powerdns_recursor.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 jmx.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 couch.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 gearmand.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 dns_check.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 cassandra.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 prometheus.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 mcache.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 ceph.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 btrfs.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul  5 19:54 disk.d
drwxr-xr-x 2 dd-agent dd-agent 4096 Jul 11 14:49 mongo.d
-rw-r--r-- 1 root      root      91 Jul 11 18:47 mycheck.yaml
vagrant@vagrant:/etc/datadog-agent/conf.d$ vi mycheck.yaml
vagrant@vagrant:/etc/datadog-agent/conf.d$ cd ..
vagrant@vagrant:/etc/datadog-agent$ cd checks.d/
vagrant@vagrant:/etc/datadog-agent/checks.d$ ls
mycheck.py  mycheck.pyc
vagrant@vagrant:/etc/datadog-agent/checks.d$

```

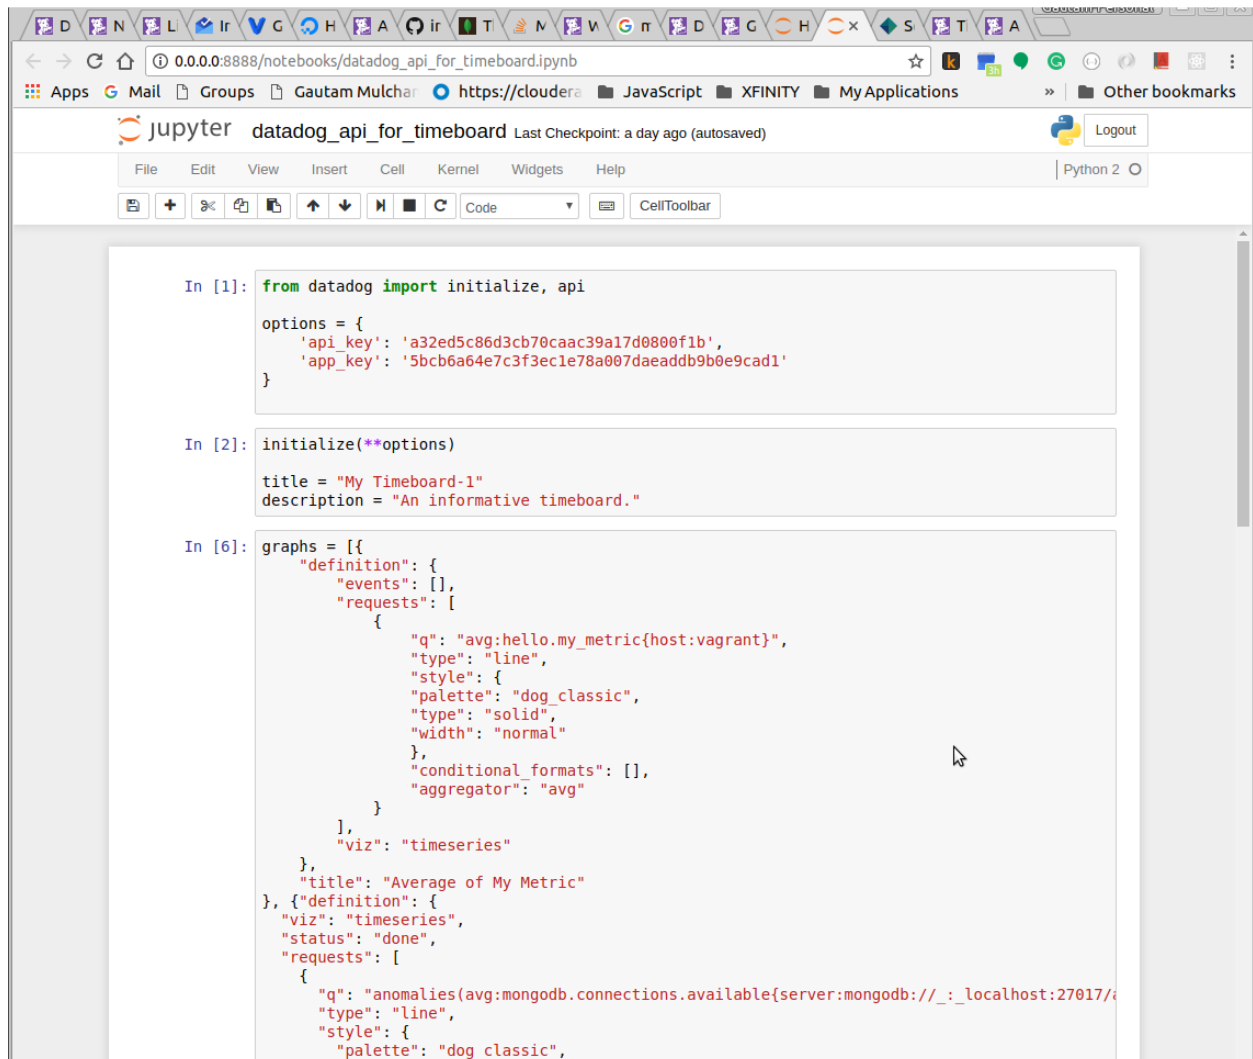


```
from checks import AgentCheck
from random import randint

class MyCheck(AgentCheck):
    def check(self, instance):
        self.gauge('hello.world', 1)
        self.gauge('hello.my_metric', randint(1,1001))
```

Visualizing Data:

Datadog dashboarding capabilities look impressive. However creating them through APIs required some reading and understand of the various parameters. So it was very nice to create a dashboard in the UI and then plug in the JSON of each graph into the API call. Testing done in Jupyter notebook and screenshots below (Code also attached in different files):



The screenshot shows a Jupyter Notebook interface with the following code cells:

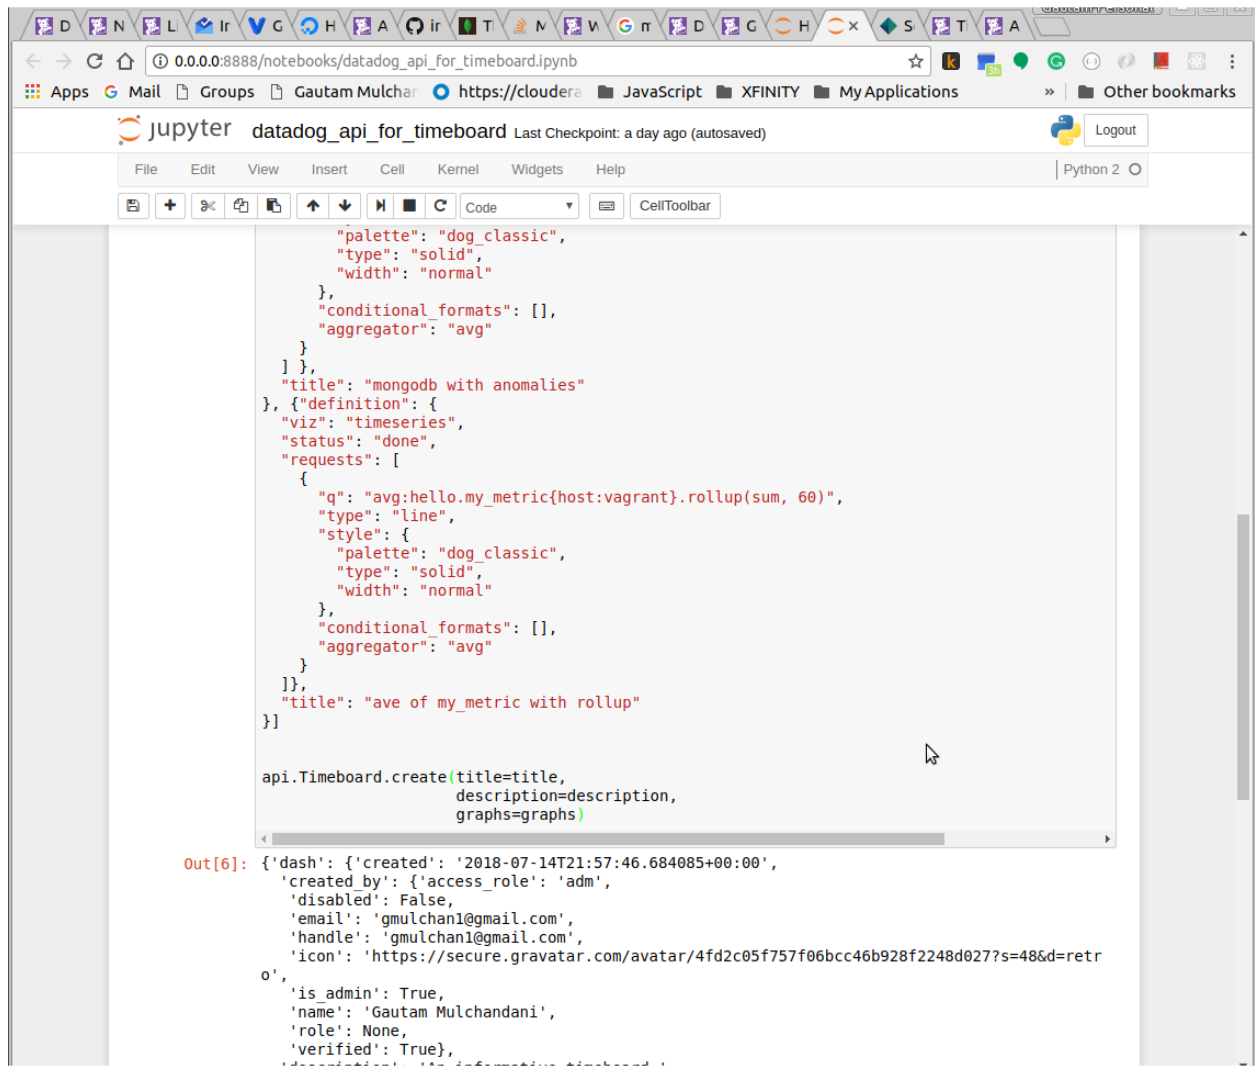
```
In [1]: from datadog import initialize, api

options = {
    'api_key': 'a32ed5c86d3cb70caac39a17d0800f1b',
    'app_key': '5bcb6a64e7c3f3ec1e78a007daeaddb9b0e9cad1'
}
```

```
In [2]: initialize(**options)

title = "My Timeboard-1"
description = "An informative timeboard."
```

```
In [6]: graphs = [{
    "definition": {
        "events": [],
        "requests": [
            {
                "q": "avg:hello.my_metric(host:vagrant)",
                "type": "line",
                "style": {
                    "palette": "dog_classic",
                    "type": "solid",
                    "width": "normal"
                },
                "conditional_formats": [],
                "aggregator": "avg"
            }
        ]
    },
    "viz": "timeseries"
}, {
    "title": "Average of My Metric"
}, {
    "definition": {
        "viz": "timeseries",
        "status": "done",
        "requests": [
            {
                "q": "anomalies(avg:mongodb.connections.available(server:mongodb://_:_localhost:27017/"))",
                "type": "line",
                "style": {
                    "palette": "dog_classic",
                    "type": "solid",
                    "width": "normal"
                },
                "conditional_formats": [],
                "aggregator": "avg"
            }
        ]
    },
    "viz": "timeseries"
}]
```

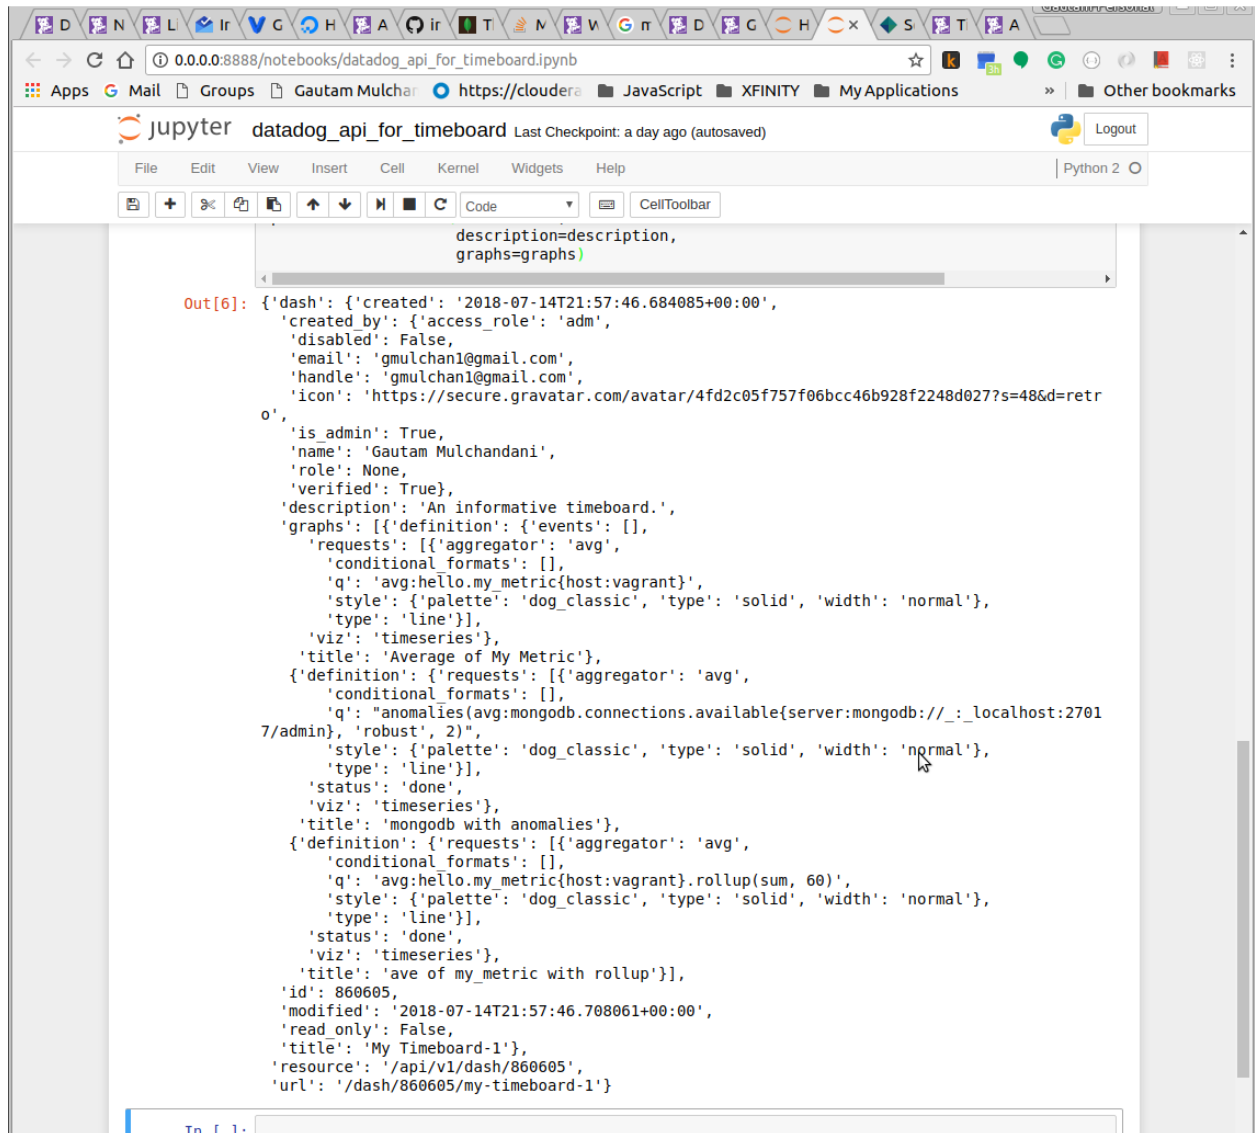


The screenshot shows a Jupyter Notebook titled "datadog_api_for_timeboard" with a last checkpoint from a day ago. The notebook is running on a local server at 0.0.0.0:8888. The code in the cell defines a timeboard with two graphs: "mongodb with anomalies" and "ave of my_metric with rollup". The output shows the successful creation of the timeboard and the user's profile information.

```
    "palette": "dog_classic",
    "type": "solid",
    "width": "normal"
  },
  "conditional_formats": [],
  "aggregator": "avg"
},
{"title": "mongodb with anomalies",
 "definition": {
  "viz": "timeseries",
  "status": "done",
  "requests": [
    {
      "q": "avg:hello.my_metric{host:vagrant}.rollup(sum, 60)",
      "type": "line",
      "style": {
        "palette": "dog_classic",
        "type": "solid",
        "width": "normal"
      },
      "conditional_formats": [],
      "aggregator": "avg"
    }
  ]
},
{"title": "ave of my_metric with rollup"}
]

api.Timeboard.create(title=title,
                     description=description,
                     graphs=graphs)
```

```
Out[6]: {'dash': {'created': '2018-07-14T21:57:46.684085+00:00',
                  'created_by': {'access_role': 'adm',
                                'disabled': False,
                                'email': 'gmulchan1@gmail.com',
                                'handle': 'gmulchan1@gmail.com',
                                'icon': 'https://secure.gravatar.com/avatar/4fd2c05f757f06bcc46b928f2248d027?s=48&d=retr
o',
                                'is_admin': True,
                                'name': 'Gautam Mulchandani',
                                'role': None,
                                'verified': True},
                  'name': 'Timeboard',
                  'status': 'active',
                  'updated': '2018-07-14T21:57:46.684085+00:00',
                  'updated_by': {'access_role': 'adm',
                                'disabled': False,
                                'email': 'gmulchan1@gmail.com',
                                'handle': 'gmulchan1@gmail.com',
                                'icon': 'https://secure.gravatar.com/avatar/4fd2c05f757f06bcc46b928f2248d027?s=48&d=retr
o',
                                'is_admin': True,
                                'name': 'Gautam Mulchandani',
                                'role': None,
                                'verified': True},
                  'version': '1.0.0'}}
```

```
description=description,
graphs=graphs)

Out[6]: {'dash': {'created': '2018-07-14T21:57:46.684085+00:00',
'created_by': {'access_role': 'adm',
'disabled': False,
'email': 'gmulchan1@gmail.com',
'handle': 'gmulchan1@gmail.com',
'icon': 'https://secure.gravatar.com/avatar/4fd2c05f757f06bcc46b928f2248d027?s=48&d=retr
o',
'is_admin': True,
'name': 'Gautam Mulchandani',
'role': None,
'verified': True},
'description': 'An informative timeboard.',
'graphs': [{'definition': {'events': [],
'requests': [{'aggregator': 'avg',
'conditional_formats': [],
'q': 'avg:hello.my_metric{host:vagrant}',
'style': {'palette': 'dog_classic', 'type': 'solid', 'width': 'normal'},
'type': 'line'}],
'viz': 'timeseries'},
'title': 'Average of My Metric'},
{'definition': {'requests': [{'aggregator': 'avg',
'conditional_formats': [],
'q': 'anomalies(avg:mongodb.connections.available{server:mongodb://_:localhost:2701
7/admin}, 'robust', 2)',
'style': {'palette': 'dog_classic', 'type': 'solid', 'width': 'normal'},
'type': 'line'}],
'status': 'done',
'viz': 'timeseries'},
'title': 'mongodb with anomalies'},
{'definition': {'requests': [{'aggregator': 'avg',
'conditional_formats': [],
'q': 'avg:hello.my_metric{host:vagrant}.rollup(sum, 60)',
'style': {'palette': 'dog_classic', 'type': 'solid', 'width': 'normal'},
'type': 'line'}],
'status': 'done',
'viz': 'timeseries'},
'title': 'ave of my_metric with rollup'}],
'id': 860605,
'modified': '2018-07-14T21:57:46.708061+00:00',
'read_only': False,
'title': 'My Timeboard-1',
'resource': '/api/v1/dash/860605',
'url': '/dash/860605/my-timeboard-1'}
```

Below, the My Timeboard-1 is created using datadog API and Gautam's Timeboard was created directly in UI.

Secure | <https://app.datadoghq.com/dashboard/lists>

Apps | Mail | Groups | Gautam Mulchar | <https://cloudere> | JavaScript | XFINITY | My Applications | Other bookmarks

Welcome, Gautam! | Get started | You are 67% done setting up. | You have 6 days left in your trial. Upgrade

Dashboards

New List + | New Dashboard +

Search dashboards







Lists

Filter lists

Preset Lists

- All Custom
- All Hosts
- All Integrations
- Created By You
- Frequently Viewed By You

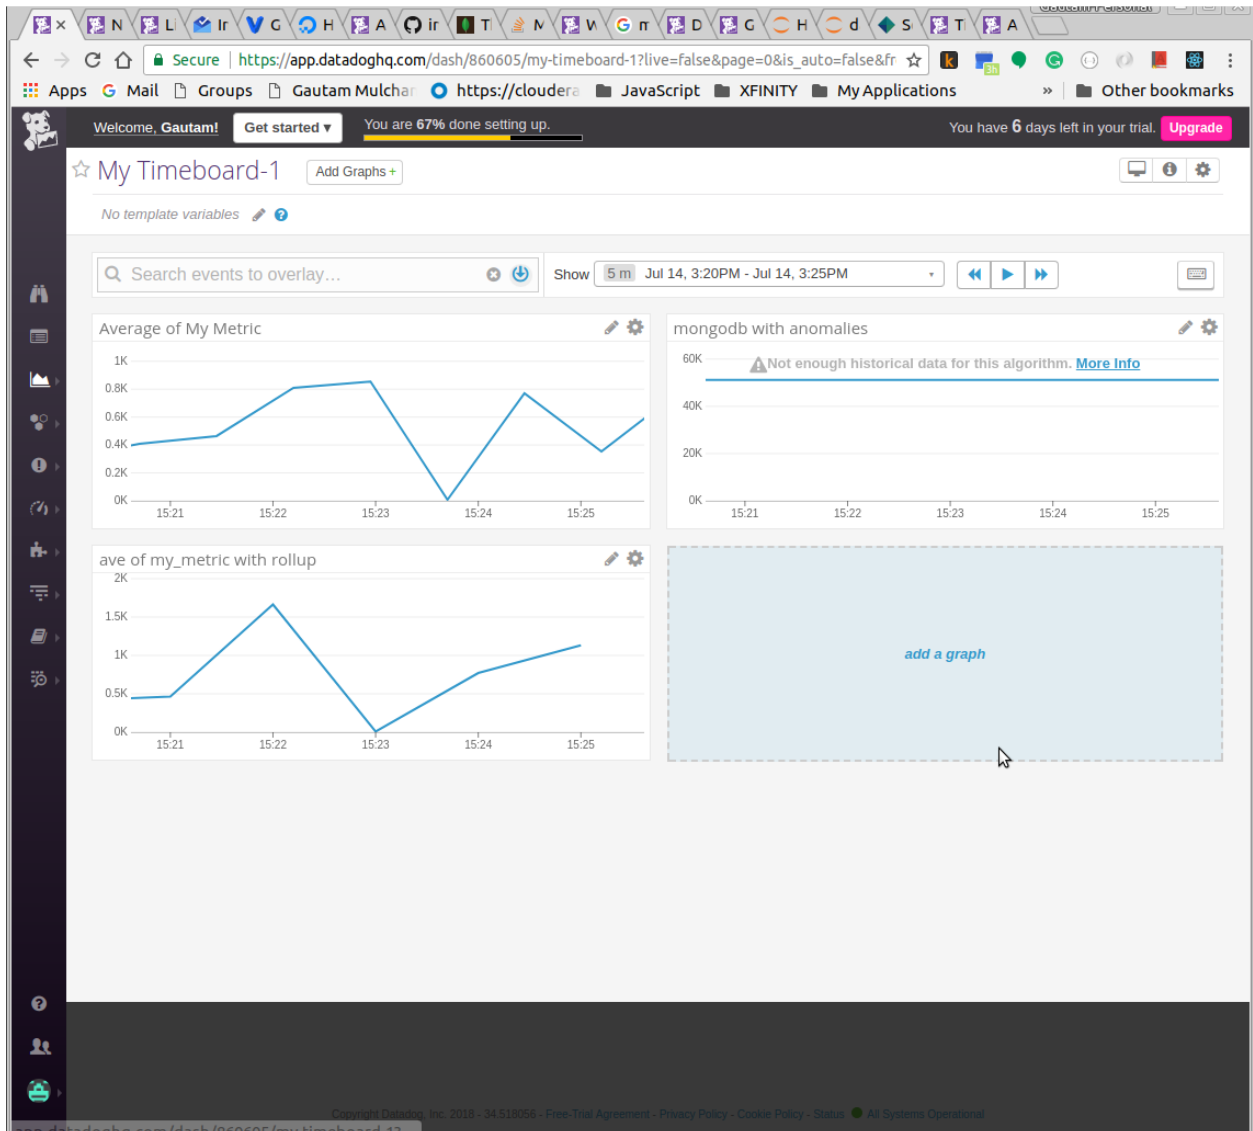
All Dashboards

<input type="checkbox"/>	Name	Modified ↑	
<input type="checkbox"/>	My Timeboard-1	Jul 14, 2018	
<input type="checkbox"/>	Gautam's Timeboard 13 Jul 2018 15:23	Jul 13, 2018	
<input type="checkbox"/>	MongoDB <small>Preset</small>	-	
<input type="checkbox"/>	System - Disk I/O <small>Preset</small>	-	
<input type="checkbox"/>	System - Networking <small>Preset</small>	-	
<input type="checkbox"/>	System - Overview <small>Preset</small>	-	

Copyright Datadog, Inc. 2018 - 34.518056 - [Free-Trial Agreement](#) - [Privacy Policy](#) - [Cookie Policy](#) - [Status](#) | All Systems Operational

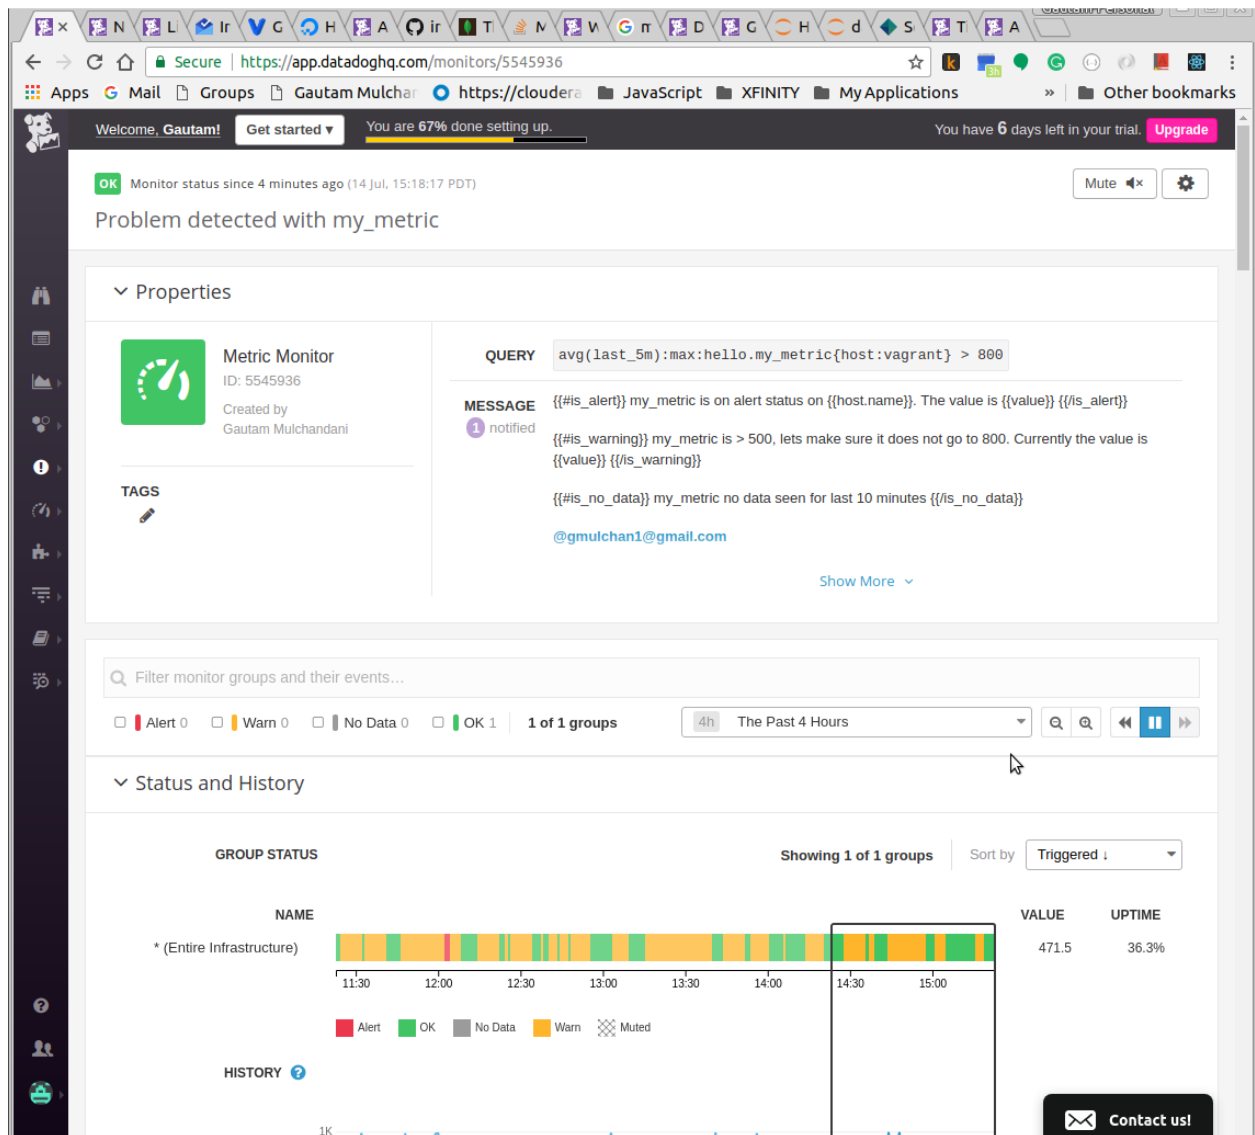
Contact us!

Last 5 minutes on graph:



Can send a graph in email by annotating it with @email address but not the entire timeboard. The anomaly graph just shows not enough data - probably because it needs sufficient historical data points to compare to and make predictions on what is “normal”.

Monitoring Data.



The alert from the monitor that was setup is below. Setting up monitor and alert quite simple, even with different messages for warning, critical and no data.

[Monitor Alert] Warn: Problem detected with my_metric



Datadog Alerting to me


10:04 AM



[Warn] Problem detected with my_metric

my_metric is > 500, lets make sure it does not go to 800. Currently the value is 553.833

[@gmulchan1@gmail.com](mailto:gmulchan1@gmail.com)



hello.my_metric over host:vagrant was > 500.0 on average during the last 5m.

The monitor was last triggered at Sat Jul 14 2018 17:04:16 UTC (1 sec ago).

[\[Monitor Status\]](#) · [\[Edit Monitor\]](#) · [\[Show Processes\]](#) · [\[Related Logs\]](#)

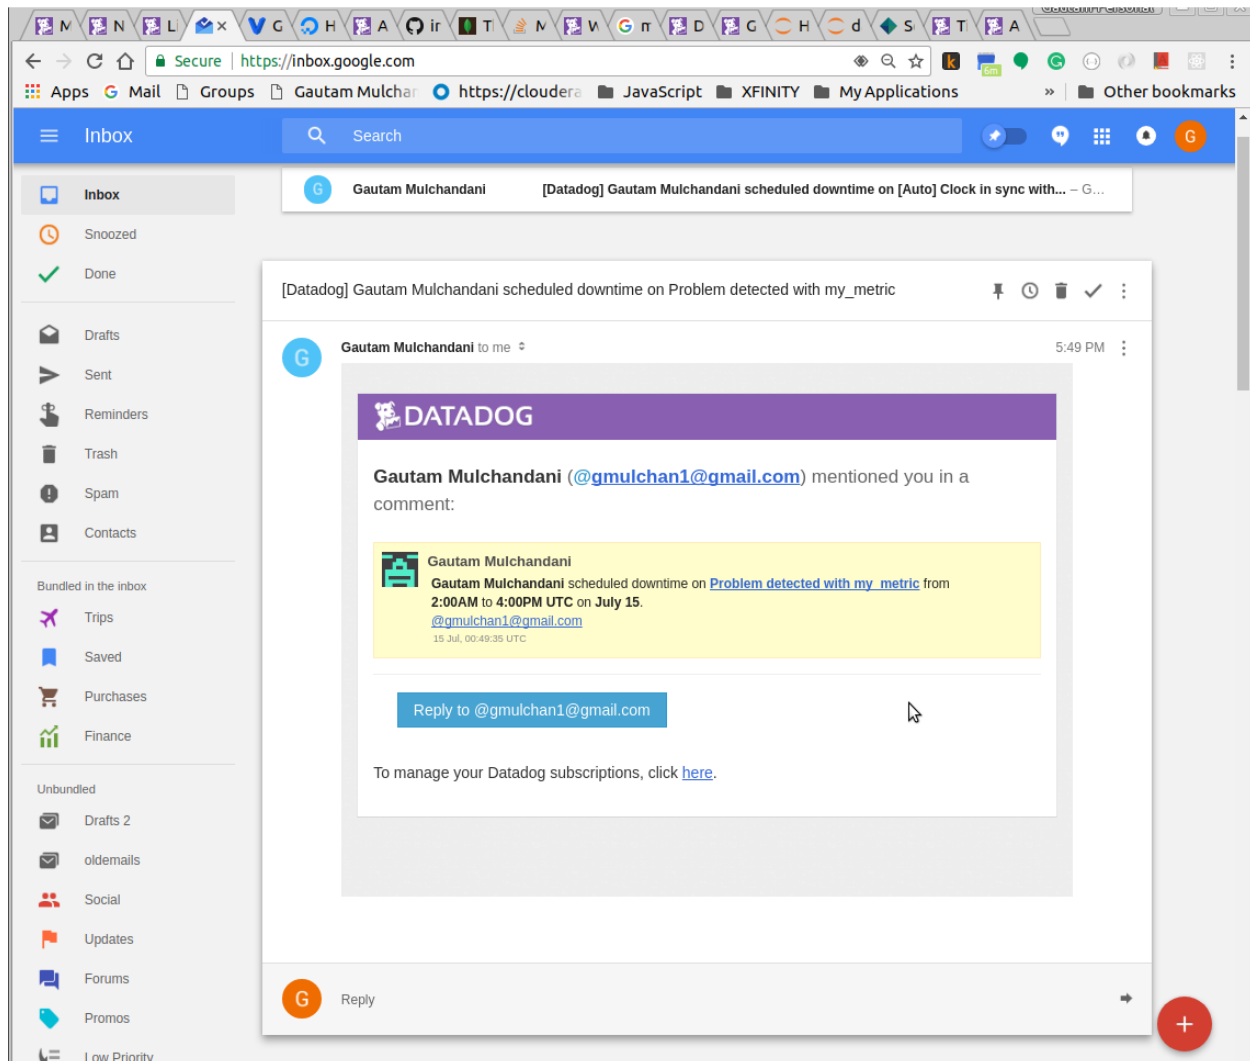
This alert was raised by account Datadog Recruiting Candidate

[Comment in Datadog](#)

To manage your Datadog subscriptions, click [here](#).

Scheduling downtime:

Important to provide a way to schedule downtime. Downtime can occur as a result of regular maintenance, for example during upgrade or changing hardware resources. If downtime is not scheduled, it can result in an alert storm where the alerts are meaningless. Also, does not help the operations team in reporting their SLAs.



Collecting APM Data

Combining both APM and infrastructure monitoring into single dashboards is very powerful as one can quickly make time based correlations between the Business Services that are served by the Applications and the infrastructure components the applications rely upon. This helps in triaging and even diagnosing problems instead of multiple teams in development, operations, networks, databases pointing fingers at each other. Adding APM metrics to the existing dashboard was simple by just going to the APM screen and selecting Export to Timeboard and then selecting the Timeboard.

Also liked that there is no need for code changes. The datadog agent handles the instrumentation automatically. However, to add traces have to put the annotation in the

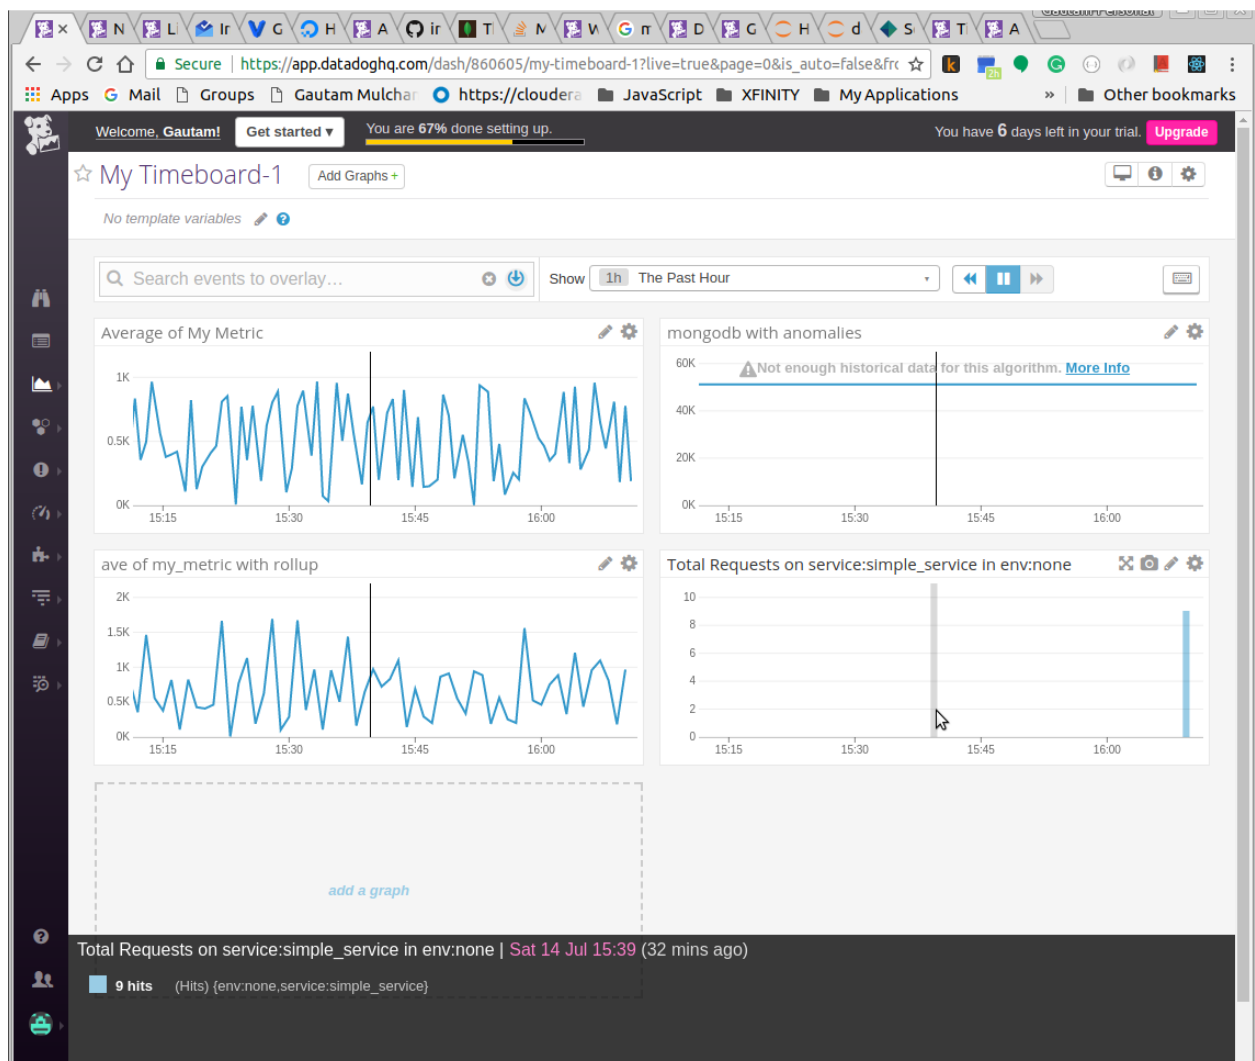
source code. There might be another way to achieve the same through monkey patching.

The modified source code for the Flask application is also uploaded.

Link to Timeboard with both APM and Infra metrics.

https://app.datadoghq.com/dash/860605/my-timeboard-1?live=true&page=0&is_auto=false&from_ts=1531606346480&to_ts=1531609946480&tile_size=m

Screenshot:



Difference between Service and Resource:

<https://help.datadoghq.com/hc/en-us/articles/115000702546-What-is-the-Difference-Between-Type-Service-Resource-and-Name->

At a high level, a Service is a logical entity that is implemented by an application. The service itself provides resources, each resource implemented by an API call.

Is there anything creative you would use Datadog for?

Looking at capabilities of datadog while performing the exercises, I would consider using it for predicting scenarios and taking action that are beneficial.

1. For example, Mumbai city gets flooded during monsoon season as it happened in the last 1 week. Based on the amount of rainfall, water level of certain streets, level of the sea tide, one can predict exactly which neighbourhood and streets will get flooded next and take preventive measures such as diverting traffic, pumping water out at certain intersections. This can bring relief to people who have to guess while navigating the heavy rainfall and still get to work.
2. Measure the temperature, general weather patterns, vaccinations given to population and use that as a measure of how many patients will be admitted to hospitals. That will determine how many personnel are required - nurses, doctors as well as medications. Proactively have the resources ready or alerted so that the right amount of resources are available to people who need help.